

# Taking Omid to the Clouds:

## Fast, Scalable Transactions for Real-Time Cloud Analytics

O. Shacham • Y. Gottesman • A. Bergman • E. Bortnikov • E. Hillel • I. Keidar  
Yahoo Research

### Omid

**Transactional API** over **NoSQL** key value  
Client Library + Runtime Service  
Open source – **Apache incubator**  
Implemented over **Apache HBase**  
**Snapshot Isolation** consistency  
Highly Available  
Originally optimized for **throughput**

### Omid Design Choices

System	Validation	Commit entry writes	Multi tenancy
Perculator (2PC)	D	D	no
CockroachDB	D	D	yes
Omid	C	C	yes
Omid LL	C	D	yes

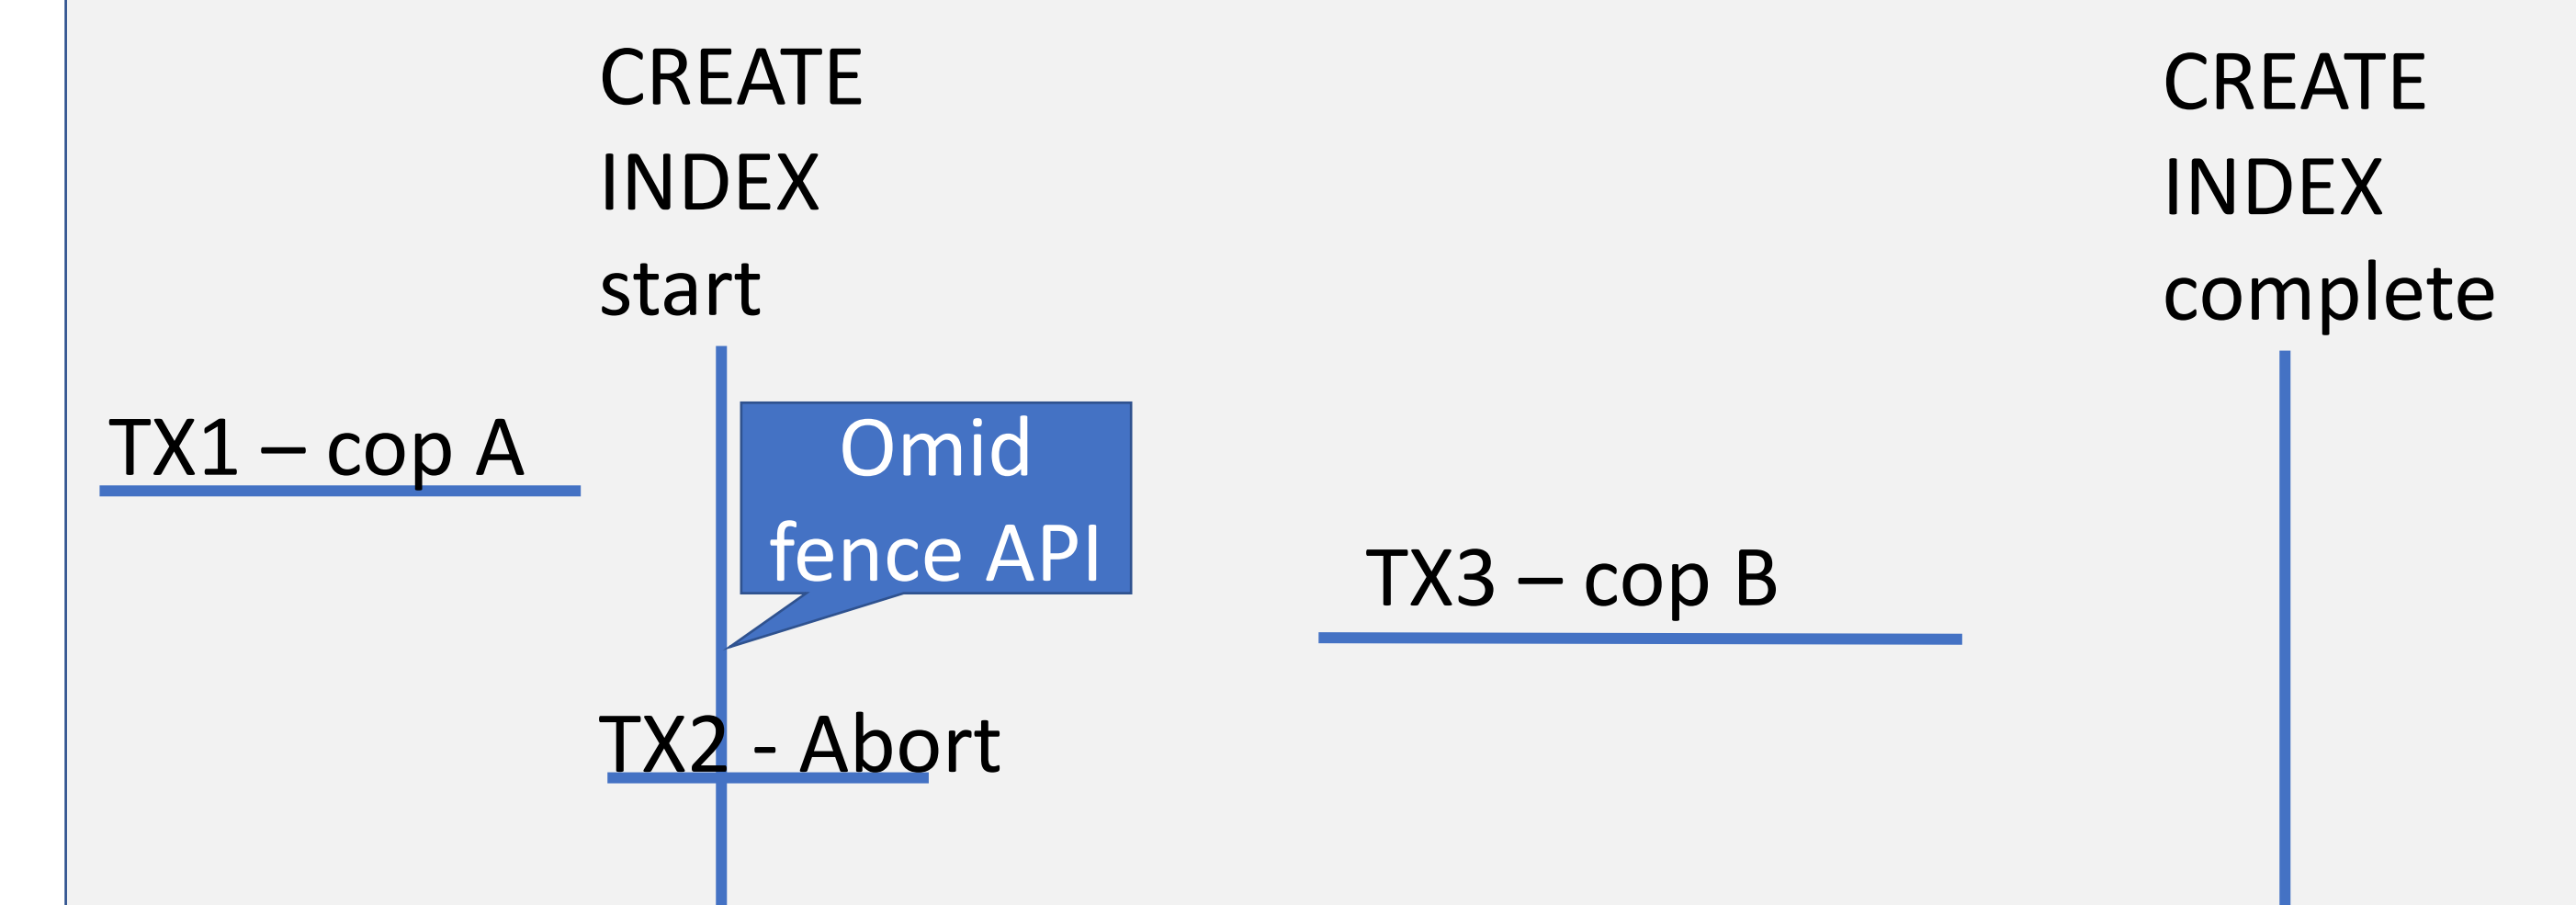
C – centralized, D – distributed

### Phoenix Integration

#### Secondary index creation

When a new index is created, two coprocessors are installed:

- Populates new index with history
- Maintains index by augmenting write to table with a write to index



### Taking Omid To The Cloud

#### Optimized for low latency – Omid LL

Redesign Omid to eliminate key bottleneck  
Distribute commit table updates among clients.

#### Novel Fast Path API – Omid FP

Many workloads have single key transactions  
Eliminate transaction overhead for them


Direct access to region server without access to TM

API:

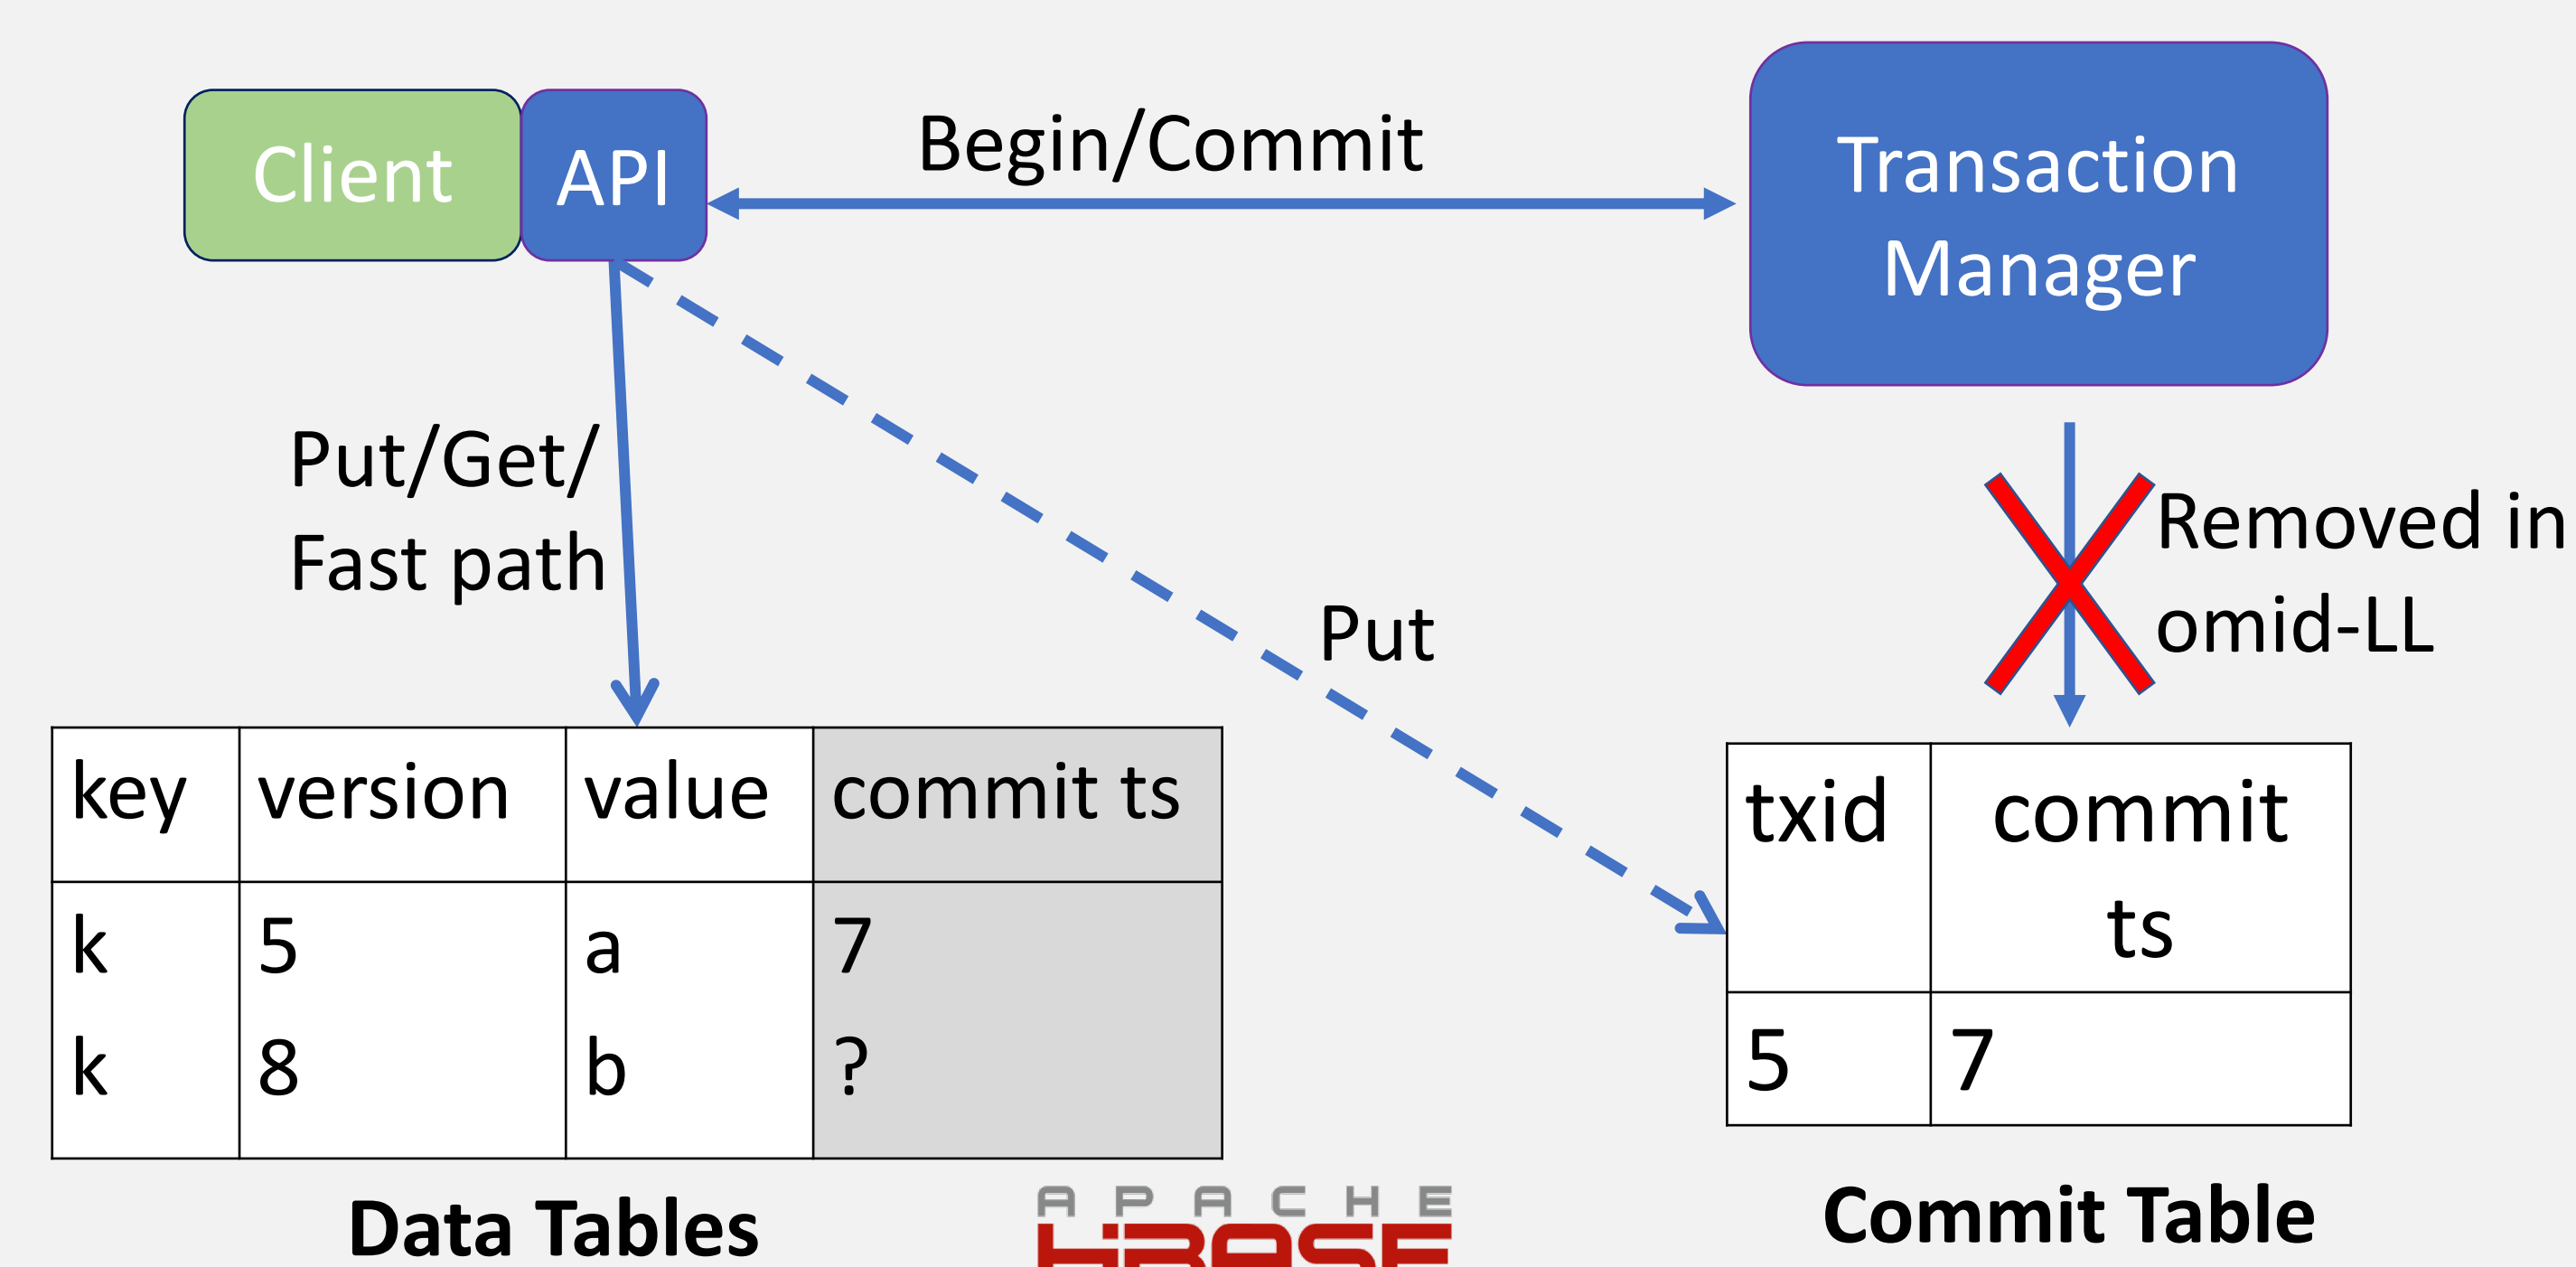
- $brc(key)$
- $bwc(key, value)$
- $br(key) + wc(key, val)$

#### Apache Phoenix Integration

Integrate Omid as the transactional layer in Phoenix SQL engine

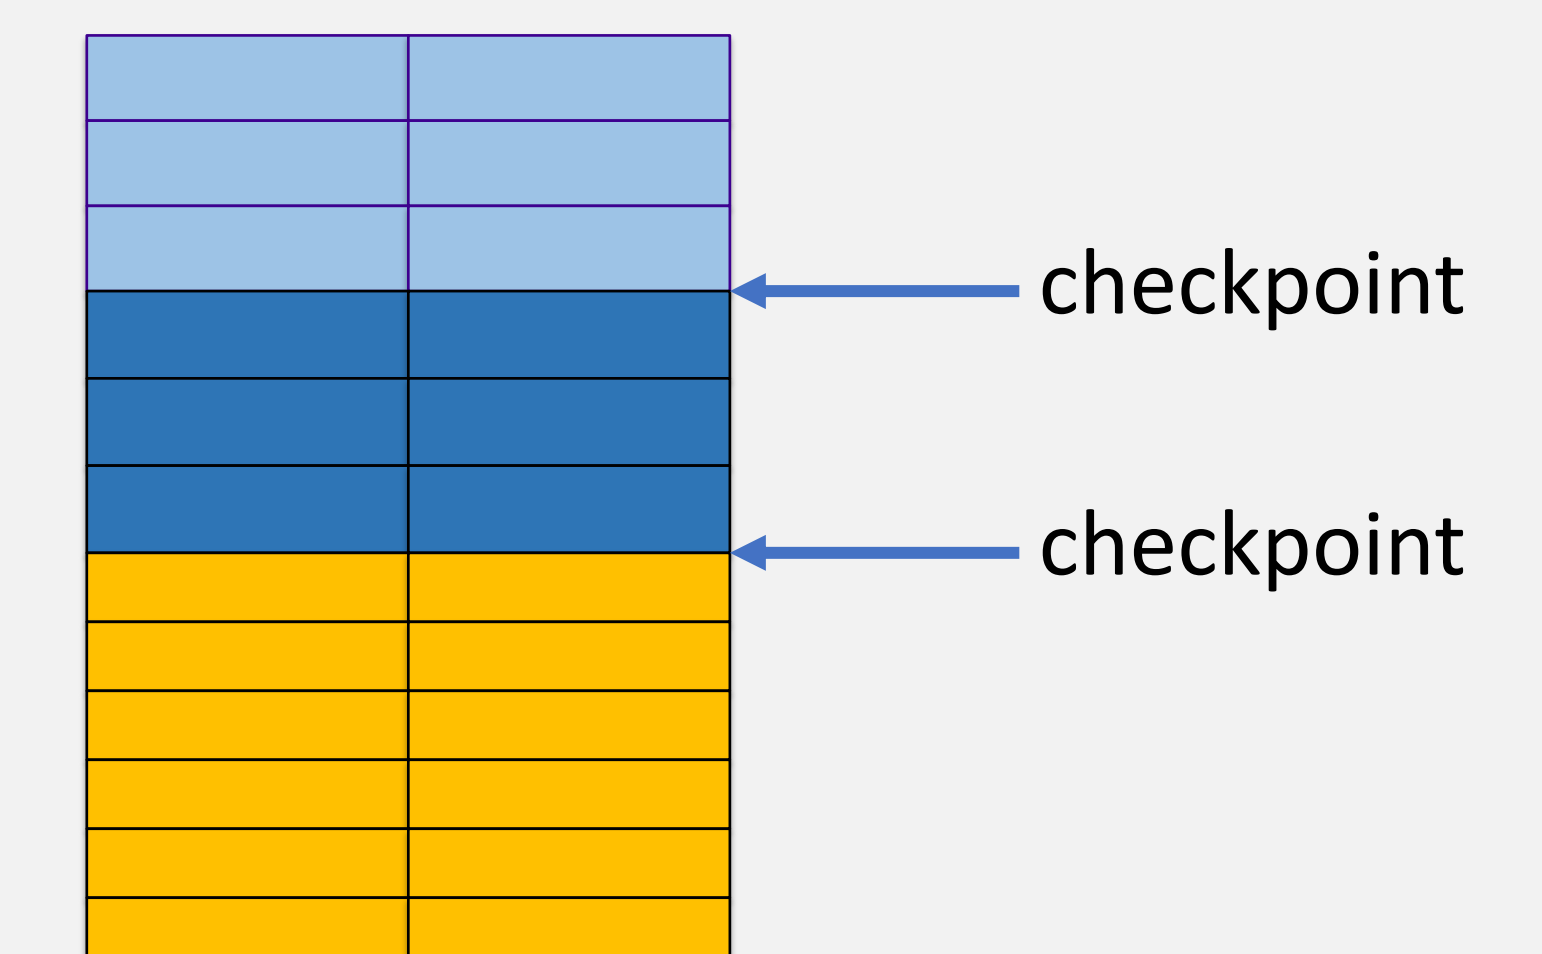


### Omid Architecture



#### Beyond snapshot isolation - SIX

INSERT INTO T  
SELECT ID+10 FROM T;  
  
INSERT INTO T  
SELECT ID+100 FROM T;



### Performance

